

HADOOP 3.2.0 MULTI NODE INSTALLATION

- sudo apt-get update && sudo apt-get -y dist-upgrade
OR
- sudo -- sh -c 'apt-get update; apt-get upgrade -y; apt-get dist-upgrade -y; apt-get autoremove -y; apt-get autoclean -y'
- sudo apt-get install openjdk-8-jdk
- java -version
- sudo addgroup hadoop
- sudo adduser --ingroup hadoop hduser
- sudo adduser hduser sudo
- sudo apt-get install ssh

After cloning:

for all nodes:

1- Change host name:

- sudo gedit /etc/hostname

Type ---> **master, etc.**

2- Change hosts:

- sudo gedit /etc/hosts

Type --->

192.168.56.10 master

192.168.56.11 slave1

192.168.56.12 slave2

192.168.56.6 dataStore

3- Make the Ip stable:

- sudo gedit /etc/network/interfaces

interfaces(5) file used by ifup(8) and ifdown(8)

auto lo

iface lo inet loopback

auto enp0s8

iface enp0s8 inet static

address 192.168.56.**10**

netmask 255.255.255.0

4- sudo reboot

5- check them ---> ifconfig

Then for SSH:

- sudo su - hduser
- sudo ufw disable
- ssh-keygen -t rsa -P "" -f ~/.ssh/id_rsa
- cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
- chmod 700 ~/.ssh/id_rsa
- ssh-copy-id -i /home/hduser/.ssh/id_rsa.pub hduser@192.168.56.10
- ssh-copy-id -i /home/hduser/.ssh/id_rsa.pub hduser@192.168.56.11
- ssh-copy-id -i /home/hduser/.ssh/id_rsa.pub hduser@192.168.56.12
- ssh-copy-id -i /home/hduser/.ssh/id_rsa.pub hduser@192.168.56.13
- ssh-copy-id -i /home/hduser/.ssh/id_rsa.pub umit@192.168.56.6

Then try to ping each other ---> ping 192.168.56.6

And try SSH ---> ssh hduser@192.168.56.6

OR

ssh hduser@dataStore

Download Hadoop

- wget <https://archive.apache.org/dist/hadoop/core/hadoop-3.2.0/hadoop-3.2.0.tar.gz> -P ~/Downloads
- wget https://archive.apache.org/dist/hadoop/common/hadoop-3.2.1/hadoop-3.2.1.tar.gz -P ~/Downloads
- sudo tar zxvf ~/Downloads/hadoop-* -C /usr/local
- sudo mv /usr/local/hadoop-* /usr/local/hadoop
- sudo chown -R hduser:hadoop /usr/local/hadoop
- sudo gedit ~/.bashrc

```
-----
export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64
export PATH=$PATH:$JAVA_HOME/bin
export HADOOP_HOME=/usr/local/hadoop
export PATH=$PATH:$HADOOP_HOME/bin
export HADOOP_CONF_DIR=/usr/local/hadoop/etc/hadoop
export HADOOP_MAPRED_HOME=${HADOOP_HOME}
export HADOOP_COMMON_HOME=${HADOOP_HOME}
export HADOOP_HDFS_HOME=${HADOOP_HOME}
export YARN_HOME=${HADOOP_HOME}
```

- ```

- source ~/.bashrc
```

### on ALL HADOOP NODES

- cd \$HADOOP\_CONF\_DIR
- sudo gedit hadoop-env.sh

Replace this line:

```
export JAVA_HOME=${JAVA_HOME}
with the following line:
export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64
```

### on ALL HADOOP NODES

- **sudo gedit core-site.xml**

```
<configuration>
<property>
 <name>fs.defaultFS</name>
 <value>hdfs://master:9000</value>
</property>
</configuration>
```

### ONLY namenode: Create namenode dir

- `sudo rm -rf /usr/local/hadoop/hdfs/namenodedata`
- `sudo mkdir -p /usr/local/hadoop/hdfs/namenodedata`
- `sudo chown -R hduser:hadoop /usr/local/hadoop/hdfs/namenodedata`
- `sudo chmod 700 -R '/usr/local/hadoop/hdfs/namenodedata'`

### All Nodes: Create Data dir (do this on namenode if it is also a datanode. Otherwise do this for only datanodes)

- `sudo rm -rf /usr/local/hadoop/hdfs/datanodedata`
- `sudo mkdir -p /usr/local/hadoop/hdfs/datanodedata`
- `sudo chown -R hduser:hadoop /usr/local/hadoop/hdfs/datanodedata`
- `sudo chmod 700 -R '/usr/local/hadoop/hdfs/datanodedata'`

### Namenode: Setup HDFS Properties

- **sudo gedit hdfs-site.xml**

```
<configuration>
<property>
 <name>dfs.replication</name>
 <value>3</value>
</property>
<property>
 <name>dfs.namenode.name.dir</name>
 <value>file:///usr/local/hadoop/hdfs/namenodedata</value>
</property>
<property>
 <name>dfs.datanode.data.dir</name>
```

```
<value>file:///usr/local/hadoop/hdfs/datanodedata</value>
</property>
</configuration>
```

### **Namenode: Setup MapReduce Properties ( with JobHistoryServer )**

- **sudo gedit mapred-site.xml**

```
<configuration>
<property>
 <name>mapreduce.jobtracker.address</name>
 <value>master:54311</value>
</property>
<property>
 <name>mapreduce.framework.name</name>
 <value>yarn</value>
</property>
<property>
 <name>mapreduce.jobhistory.address</name>
 <value>master:10020</value>
</property>
<property>
 <name>mapreduce.jobhistory.webapp.address</name>
 <value>master:19888</value>
</property>
<property>
 <name>yarn.app.mapreduce.am.env</name>
 <value>HADOOP_MAPRED_HOME=$HADOOP_HOME</value>
</property>
<property>
 <name>mapreduce.map.env</name>
 <value>HADOOP_MAPRED_HOME=$HADOOP_HOME</value>
</property>
<property>
 <name>mapreduce.map.speculative</name>
 <value>true</value>
</property>
<property>
 <name>mapreduce.reduce.speculative</name>
 <value>>false</value>
</property>
<property>
 <name>mapreduce.reduce.env</name>
 <value>HADOOP_MAPRED_HOME=$HADOOP_HOME</value>
</property>
<property>
 <name>mapreduce.job.reduce.slowstart.completedmaps</name>
 <value>0.99</value>
```

```
</property>
</configuration>
```

### ALL: Setup YARN Properties

- **sudo gedit yarn-site.xml**

```
<configuration>
 <property>
 <name>yarn.nodemanager.aux-services</name>
 <value>mapreduce_shuffle</value>
 </property>
 <property>
 <name>yarn.nodemanager.aux-services.mapreduce_shuffle.class</name>
 <value>org.apache.hadoop.mapred.ShuffleHandler</value>
 </property>
 <property>
 <name>yarn.resourcemanager.hostname</name>
 <value>master</value>
 </property>
 <property>
 <name>yarn.resourcemanager.resource-tracker.address</name>
 <value>master:8025</value>
 </property>
 <property>
 <name>yarn.resourcemanager.scheduler.address</name>
 <value>master:8030</value>
 </property>
 <property>
 <name>yarn.resourcemanager.address</name>
 <value>master:8050</value>
 </property>
 <property>
 <name>yarn.log-aggregation-enable</name>
 <value>true</value>
 </property>
 <property>
 <name>yarn.resourcemanager.scheduler.class</name>
 <value>org.apache.hadoop.yarn.server.resourcemanager.scheduler.fair.FairScheduler</value>
 </property>
</configuration>
```

### Old one:

```
<configuration>
<!-- Site specific YARN configuration properties -->
<property>
<name>yarn.nodemanager.aux-services</name>
<value>mapreduce_shuffle</value>
```

```
</property>
<property>
<name>yarn.resourcemanager.hostname</name>
<value>namenode</value>
</property>
</configuration>
```

### **Namenode: Setup Master and Slaves**

- `cd $HADOOP_CONF_DIR`
- `sudo cp workers masters`
- `sudo gedit masters`

**Delete localhost and add only the line below:**

master

- `sudo gedit workers`

**Delete localhost and add only the line below: (add only datanode's Private IPs)**

master

slave1

slave2

**Configuring Data Nodes (SLAVES) only** (don't include namenode even if it is also a datanode, as we did this above)

### **datanode: Setup HDFS Properties**

- `sudo gedit hdfs-site.xml`

```
<configuration>
<property>
 <name>dfs.replication</name>
 <value>3</value>
</property>
<property>
 <name>dfs.datanode.data.dir</name>
 <value>file:///usr/local/hadoop/hdfs/datanodedata</value>
</property>
</configuration>
```

### **Starting the Cluster**

on namenode:

- \$HADOOP\_HOME/bin/hdfs namenode -format
- \$HADOOP\_HOME/sbin/start-dfs.sh
- \$HADOOP\_HOME/sbin/start-yarn.sh
- \$HADOOP\_HOME/sbin/mr-jobhistory-daemon.sh start historyserver
- hdfs dfsadmin -safemode leave
- hdfs dfsadmin -report

### **STARTING:**

sudo su - hduser ----> qwe

sudo ufw disable

\$HADOOP\_HOME/sbin/start-dfs.sh && \$HADOOP\_HOME/sbin/start-yarn.sh &&  
 \$HADOOP\_HOME/sbin/mr-jobhistory-daemon.sh start historyserver && hdfs dfsadmin -safemode  
 leave && hdfs dfsadmin -report

### **STOPPING:**

\$HADOOP\_HOME/sbin/stop-dfs.sh && \$HADOOP\_HOME/sbin/stop-yarn.sh &&  
 \$HADOOP\_HOME/sbin/mr-jobhistory-daemon.sh stop historyserver

### **STOP:**

- \$HADOOP\_HOME/sbin/stop-dfs.sh
- \$HADOOP\_HOME/sbin/stop-yarn.sh

### **Is problem-----**

hadoop fs -mkdir /user  
 hadoop fs -mkdir /user/hduser  
 hadoop fs -ls